Module 3 Project Rubric

RiskLens | Fine-Tuned Analyst

C3 - Model Adaptation and Evaluation
Curate and document teacher-generated training data, train and reload a QLoRA Analyst, and evaluate it against the untuned base model on separated data using format, content and pairwise comparisons, with a controlled ablation and honest limitations.

Evidence to prepare

- **D3.1:** At least 150 raw input/report pairs generated by a named real teacher model, with varied products, defects and phrasing; inputs follow make_input().
- **D3.2:** Complete DATASET_CARD; deduplication and decontamination shown as executed actions with raw, removed and retained counts.
- **D3.3:** At least 25 frozen analyst_test_rows and the maximum train/test cosine-similarity check at or below 0.90, measured after decontamination.
- **D3.4:** Completed QLoRA training with configuration and loss evidence; saved, accessible and reloadable adapter; train-or-load delivery path.
- **D3.5:** Base/tuned comparisons for format adherence, content rubric score and blind randomized pairwise win-rate; named judge and caveats.
- **D3.6:** Saved Stage 3 evaluation and complete WEEK3_RESULTS, including judge identity and win/loss/tie counts, with ledger output.
- **D3.7:** One training variable changed, a second adapter trained and both runs compared on the same frozen rows; conclusion and concrete out-of-distribution failure prediction.

How grading works

All 5 criteria must be **Met** to complete C3. **C3.2 is MUST PASS** and is reviewed first. If it is Not Yet, revise before the remaining criteria are reviewed. No points, weights or bonus credit.

Submit the same notebook link with saved outputs, a milestone snapshot, a short summary and known issues. Revise Not Yet criteria and affected evidence. The full RiskLens Grading Rubric defines review rules; course policy governs attempts and timing.

Module 3 | Assessment criteria (1/2)

C3 - Model Adaptation and Evaluation

RiskLens Criterion Set v2.0 | All 5 criteria must be Met. Standards are identical to the consolidated rubric.

Criterion	Where your evidence lives	What earns "Met"	Common mistakes
C3.1 Your training data is real and curated	D3.1 teacher calls and examples; D3.2 DATASET_CARD and cleaning counts.	Generate at least 150 raw pairs with a named real teacher model and varied inputs using make_input(). Document the generation strategy. Perform deduplication and train/test decontamination, printing before/after and removal counts. Complete every dataset-card field and report retained training size. Template-only generated reports do not meet the requirement.	Using the practice template as teacher output; claiming 150 retained examples when only raw count was measured; describing cleaning without running it; incomplete card.
C3.2 Your evaluation data is separated from training MUST PASS	D3.3 analyst_test_rows; similarity calculation; dataset-card snapshot before training.	Freeze at least 25 Analyst test rows before training. Drop contaminated training rows and show the maximum train-to-test cosine similarity is at most 0.90 using a declared embedding model and method. Preserve the frozen version and counts. Explain that the similarity screen is a leakage heuristic, not proof of independence.	Training on evaluation examples; removing inconvenient test rows; reporting the mean instead of maximum similarity; using cell order as the only chronology evidence.
C3.3 Your QLoRA adapter is trained and reloadable	D3.4 training configuration, loss outputs, adapter artifact and reload demonstration.	Complete QLoRA training with a configuration appropriate to the runtime. Preserve base model ID/revision, adapter settings and training parameters, and interpret the observed loss. Save the learner-trained adapter to an accessible path or Hub location and demonstrate loading it. The delivery path loads saved weights instead of retraining on every Run all.	A configuration with no completed training; an adapter path that is inaccessible; loading only the base; no loss interpretation; retraining unexpectedly during delivery.

Evidence note. Assess saved outputs and linked artifacts; graders do not need to repeat model training. A missing evaluation output does not demonstrate a result. Preserve real prompts, configurations, data versions and outputs as the work happens. Dated snapshots or version history support freezing; cell position alone does not prove chronology.

Module 3 | Assessment criteria (2/2)

C3 - Model Adaptation and Evaluation

RiskLens Criterion Set v2.0 | All 5 criteria must be Met. Standards are identical to the consolidated rubric.

Criterion	Where your evidence lives	What earns "Met"	Common mistakes
C3.4 Your base and tuned comparison is valid	D3.5 generations and judgments; D3.6 WEEK3_RESULTS, Stage 3 output and ledger.	Evaluate genuinely untuned base and tuned outputs on the same frozen rows using format adherence, content rubric scores and blind randomized A/B judgments. Name a separate, stronger judge and state its limitations. Show wins, losses, ties and denominator. Populate WEEK3_RESULTS from saved evidence. No minimum win-rate or tuning-improvement margin is imposed.	The base still has an active adapter; fixed A/B position; a percentage without counts; judge equals generator without disclosure; scores copied without supporting outputs.
C3.5 Your ablation supports a defensible conclusion	D3.7 two adapter runs, configuration difference, comparison and conclusion.	Change exactly one training variable: rank, epochs, learning rate or training-data size. Train a second adapter and compare both runs on the same frozen rows with the same evaluation protocol. Preserve both results. Explain what changed, whether tuning is justified, a specific expected out-of-distribution failure and the limits of using the frozen set for repeated comparisons.	Multiple variables changed; no retraining; comparing different evaluation rows; hiding worse results; generic claims about generalization without a concrete failure case.

Module handoff. A lower win-rate or worse tuned output does not automatically fail C3; no improvement threshold was published for this stage. A valid comparison and justified conclusion are required. Keep judge identity and count strings visible: the ledger copies only numeric WEEK3_RESULTS values.

Source mapping: D3.1-D3.7. RiskLens Criterion Set v2.0 governs criteria; the published course policy governs administrative rules.