Module 5 Project Rubric

RiskLens | Integrated RiskLens and Final Delivery

C5 - Integrated System Delivery and Defense
Integrate the learner-built inspector, evidence engine and selected Analyst into a reproducible RiskLens pipeline, evaluate its behavior and failures across products, and explain its architecture, model choice and measurement limits using the submitted evidence.

Evidence to prepare

- **D4.5:** Working analyze_product(image), using the actual learner components; saved demonstrations on at least three products, including one good and two defective products.

- **D4.6:** Deployment note naming the target runtime, likely bottleneck at 100 times traffic and one actionable monitoring signal.

- **D4.7:** Full-pipeline test on at least ten products, with end-to-end p50/p95, stage-time breakdown and a failure taxonomy.

- **Final delivery:** Clean top-to-bottom run/load; saved evaluation outputs and show_ledger(); accessible model/data artifacts and environment details; reflection of at most 300 words.

- **Explanation evidence:** A notebook walkthrough explaining an actual component, tracing one product, defending one metric and analyzing a failed experiment or unsupported claim. A separately issued exam may supply this evidence.

How grading works

All 4 criteria must be **Met** to complete C5. **C5.1 is MUST PASS** and is reviewed first. If it is Not Yet, revise before the remaining criteria are reviewed. No points, weights or bonus credit.

Submit the same notebook link with saved outputs, a milestone snapshot, a short summary and known issues. Revise Not Yet criteria and affected evidence. The full RiskLens Grading Rubric defines review rules; course policy governs attempts and timing.

Module 5 | Assessment criteria (1/2)

C5 - Integrated System Delivery and Defense

RiskLens Criterion Set v2.0 | All 4 criteria must be Met. Standards are identical to the consolidated rubric.

Criterion	Where your evidence lives	What earns "Met"	Common mistakes
C5.1 Your real components work together MUST PASS	D4.5 analyze_product(); three saved product demonstrations; selected Analyst record.	Compose the learner inspector, retriever and selected Analyst through the published contracts. Preserve the GOOD short-circuit and return note, evidence and report. Save at least three product demonstrations: one good and two defective, with traceable inputs and outputs. Use real components, not hard-coded substitutes. Label any fallback or failed-gate demonstration and its limits.	Demo stubs standing in for components; a JSON note passed where training used prose; missing good-product behavior; reports with no traceable evidence; a rejected model labeled ready.
C5.2 Your system evaluation locates failures	D4.7 batch outputs; timing table; product-level failure taxonomy.	Run the full pipeline on at least ten products and save end-to-end p50/p95 plus inspector, retrieval and Analyst timing. For every wrong or weak report, identify the responsible stage and evidence for that diagnosis; explicitly report if none are observed. Explain what the small workload cannot establish about production reliability or generalization.	Benchmarking only the Analyst; no stage timing; vague failure categories; silently skipping failures; claiming a small batch establishes production reliability.
C5.3 Your delivery is reproducible and documented	D4.6 deployment note; final run/load and show_ledger(); artifacts, environment and reflection.	Provide a clean top-to-bottom run/load with saved evaluation outputs, accessible model/data artifacts and environment/configuration details. Display the complete metrics ledger and preserve earlier comparisons separately. Document runtime choice, the bottleneck at 100 times traffic and a monitoring signal. Include a reflection of at most 300 words on trade-offs, trusted/uncertain numbers and next steps.	Missing weights or data; a notebook that depends on hidden state; outputs cleared; assuming overwritten ledger entries preserve all experiments; deployment claims without a runnable route.

Evidence note. Assess saved outputs and linked artifacts; graders do not need to repeat model training. A missing evaluation output does not demonstrate a result. Preserve real prompts, configurations, data versions and outputs as the work happens. Dated snapshots or version history support freezing; cell position alone does not prove chronology.

Module 5 | Assessment criteria (2/2)

C5 - Integrated System Delivery and Defense

RiskLens Criterion Set v2.0 | All 4 criteria must be Met. Standards are identical to the consolidated rubric.

Criterion	Where your evidence lives	What earns "Met"	Common mistakes
C5.4 You can explain the submitted system	Explanation cells in final notebook; traceable code/outputs; exam evidence if separately issued.	Explain what one actual component does, why its design fits RiskLens and what would break if a named part were removed. Trace one demonstrated product through the applicable stages. Defend one reported metric and discuss a failed experiment or unsupported claim using saved evidence. Submit a written walkthrough; separately issued exam evidence may replace equivalent parts.	Reading code without explaining behavior; a generic architecture unrelated to the notebook; no supporting measurement; avoiding a failure or limitation; substituting polish for understanding.

Module handoff. Use a passing compressed arm when available. If all compressed candidates fail, a clearly labeled feasible reference-model fallback may demonstrate integration; it does not establish successful compression. Preserve the rejected-candidate evidence and the original gate. No paid hosting or separate recording is required.

Source mapping: D4.5-D4.7 plus final delivery. RiskLens Criterion Set v2.0 governs criteria; the published course policy governs administrative rules.